



An Approach to SLURM Configuration Verification

Kyle Reinholt

View the Slides [replace QR code and link!]



https://github.com/ResearchComputing/research_storage_simplified

Introduction – Why Configuration Matters

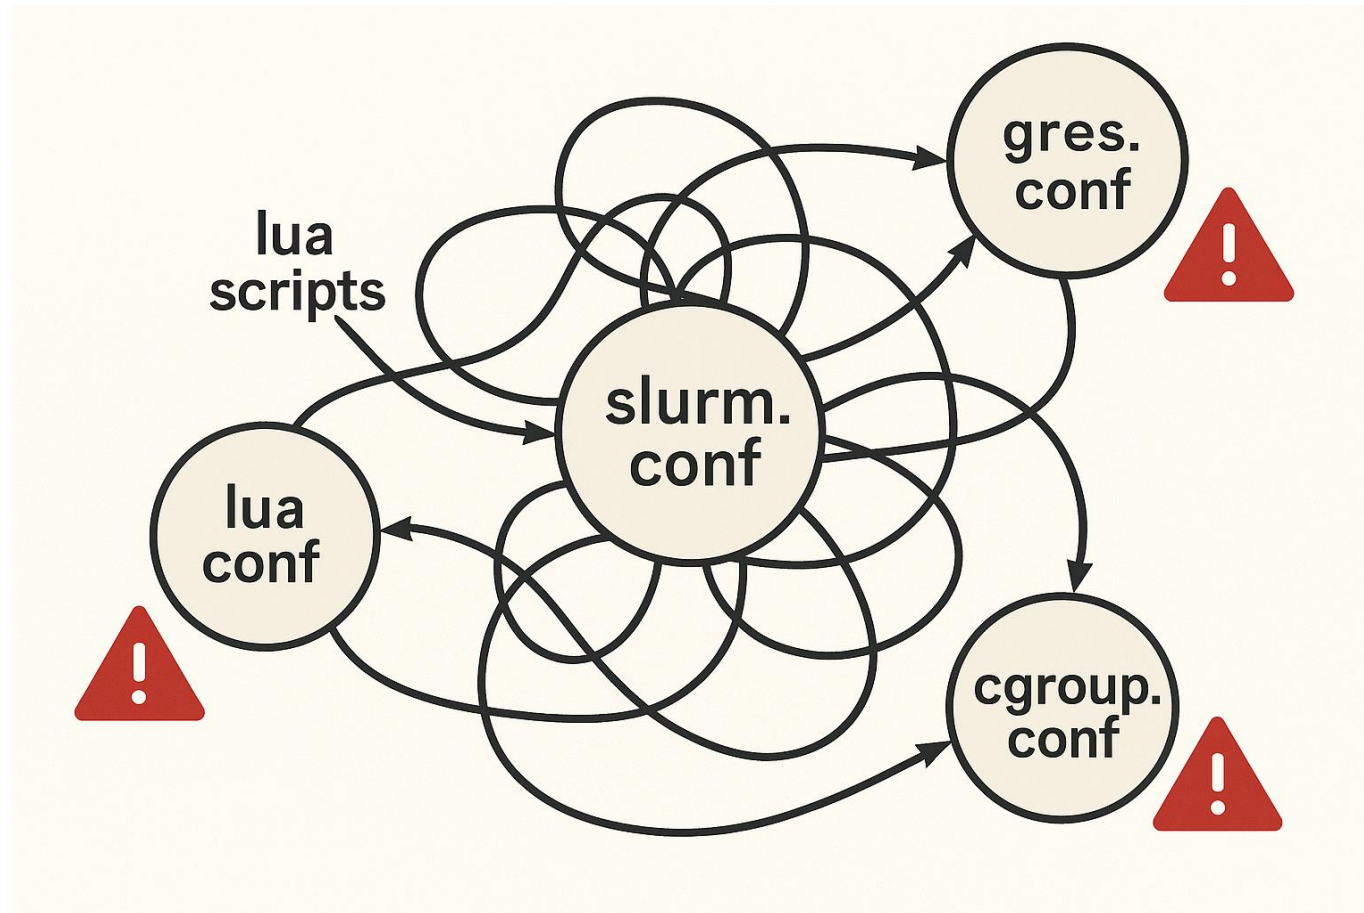
- Misconfigured SLURM settings often lead to user confusion, job failures, and a spike in support tickets — creating friction for both researchers and support staff.



The Problem

- Validating SLURM configs is painful
- Manual checks are error-prone
- You only get feedback *after* something breaks

The Problem



Current Approaches

- **Manual Reviews** — Time-consuming and error-prone
- **Cluster-Specific Scripts** — Custom, hard to reuse or share
- **Limited Tooling** — scontrol, slurmlint, UBCCR simulator offer partial checks

An Approach to SLURM Configuration Verification

- Declarative validation schemas (e.g. JSON schema for YAML-like configs)
- Automated testing pipelines
- Community-driven validation rules



Products Utilizing Schema Validation



Conclusion

- Config correctness is essential → current tools fall short → better tooling is possible
- Email - kyle@colorado.edu

Prior Art in SLURM Configuration Checking

- <https://github.com/appeltel/slurmlint> – amazing work, does not use a schema, last updated 2019
- <https://ubccr-slurm-simulator.github.io/> - incredible project for simulating configurations, supports Slurm 17.11 – is currently being updated
- Spin up slurm controller in a container and see if fails
- “If it crashes, it's not right” – Joe Malingowski

Limitations with Current Approaches

- Don't scale well
- Hard to generalize
- Not integrated into CI/CD